

# Elasticsearch

Jakub Podeszwik

Yameo

28.03.2018

- 1 Lucene
- 2 Elasticsearch
- 3 Basic data structures
- 4 Searching, Indexing
- 5 Full text search
- 6 Analyzers
- 7 Cluster

- ① Full text search library written in Java
- ② initially released in 1999
- ③ in 2010 Apache Solr joined Lucene as subproject

- ① Open source Full text search Engine written on top of lucene
- ② 02.2010 - version 0.4 released
- ③ 02.2014 - version 1.0 released
- ④ scalable, near realtime, highly available, restful

# Inverted index

1. Tom has a cat

2. Kate has a dog

3. Mike has an owl

Term	Documents
------	-----------

a	1, 2
---	------

an	3
----	---

cat	1
-----	---

dog	2
-----	---

has	1, 2, 3
-----	---------

kate	2
------	---

mike	3
------	---

owl	3
-----	---

tom	1
-----	---

# Inverted index

1.

title: The cat  
text: The cat sits

2.

title: The dog  
text: The dog stands

Term	Documents
text:cat	1
text:dog	2
text:sits	1
text:stands	2
text:the	1, 2
title:cat	1
title:dog	2
title:the	1,2

Immutable, stored on disk data structure consisting of:

- 1 inverted index
- 2 fielddata cache / doc\_values
- 3 \_source
- 4 live documents bitset (this one is not immutable)
- 5 ...

## Segment 1

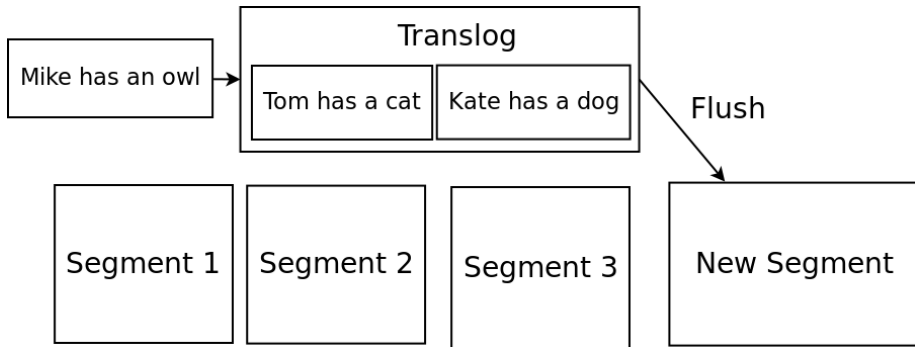
a	1, 2
cat	1
dog	2
has	1, 2
kate	2
tom	1

## Segment 2


an	3
has	3
mike	3
owl	3



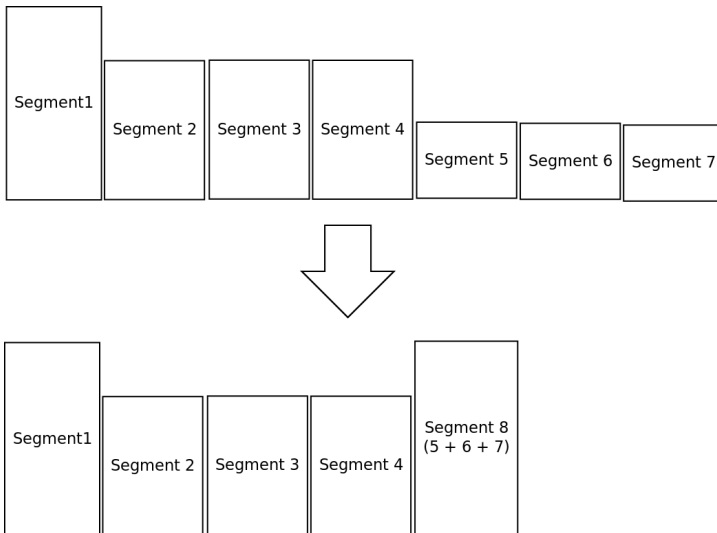
# Translog




# Segments merging




# Segments merging




# Elasticsearch index




# Elasticsearch index




# Elasticsearch shard replicas





# Search














# Search



# Search



# Indexing



- ① indices
- ② shards
- ③ nodes
- ④ master
- ⑤ ?v - legend


Definition of how documents in elasticsearch should be stored / indexed / searched.

Elasticsearch will try to guess mapping from incoming documents if some of the fields are not specified.

- 1 term
- 2 bool
- 3 prefix
- 4 wildcard
- 5 fuzzy
- 6 match, match\_all
- 7 query\_string
- 8 filter vs query



# Term



bool.should = or = sum  
{1, 2, 3} ← {1, 3}  
                  {2}

bool.should:  
term: cat  
term: dog

Term	Documents
a	1, 2
cat	1, 3
dog	2
has	1, 2
kate	2
sits	3
the	3
tom	1

1. Tom has a cat

2. Kate has a dog

3. The cat sits

bool.must = and = intersection

{1} ← {1, 3}  
          ← {1, 2}

bool.must:  
term: cat  
term: has


Term	Documents
a	1, 2
cat	1, 3
dog	2
has	1, 2
kate	2
sits	3
the	3
tom	1

1. Tom has a cat

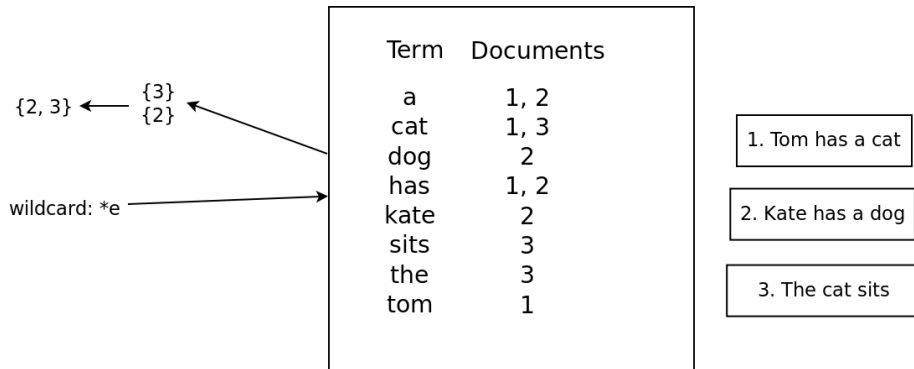
2. Kate has a dog

3. The cat sits


# Wildcard



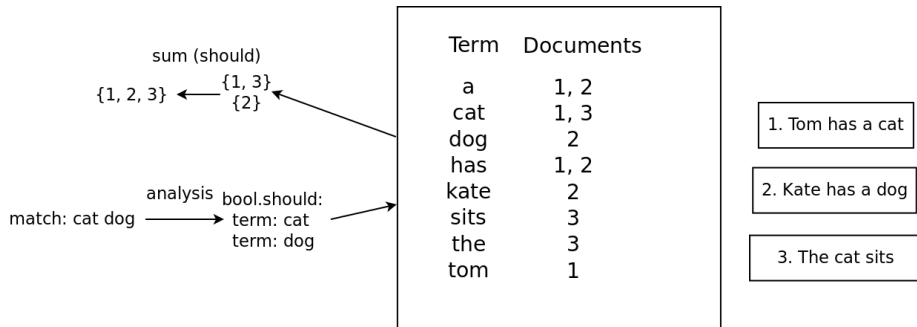
# Wildcard



# Wildcard



# Match



Analazers split input string into stream of terms. Examples:

- 1 whitespace
- 2 standard - default analyzer for text fields.
- 3 keyword - noop



# Custom analyzer

- 1 char\_filters - removes / changes / adds characters before text goes to tokenizer
- 2 tokenizer - splits text into tokens
- 3 filters - removes / changes / adds tokens

# Custom analyzer

Tom has a cat, and Kate has two dogs.

Strip punctuation char filter

Tom has a cat and Kate has two dogs

Whitespace tokenizer

[Tom, has, a, cat, and, Kate, has, two, dogs]

Lowercase filter

[tom, has, a, cat, and, kate, has, two, dogs]

English stopwords filter

[tom, cat, kate, two, dogs]

English stem filter

[tom, cat, kate, two, dog]

$$\text{score}(q, d) = \text{queryNorm}(q) * \text{coord}(q, d) * \sum (tf(t \text{ in } d) * idf(t)^2 * t.\text{getBoost()} * \text{norm}(t, d))(t \text{ in } q)$$

# Scoring

- ①  $\text{score}(q,d)$  - the relevance score of document  $d$  for query  $q$
- ②  $\text{queryNorm}(q)$  - query normalization factor
- ③  $\text{coord}(q,d)$  is the coordination factor (rewards for higher percentage of query terms contained in document)
- ④  $\text{tf}(t \text{ in } d)$  - term frequency for term  $t$  in document  $d$ ,
- ⑤  $\text{idf}(t)$  - inverse document frequency for term  $t$ ,
- ⑥  $t.\text{getBoost}()$  - boost that has been applied to the query,
- ⑦  $\text{norm}(t,d)$  - field-length norm, combined with the index-time field-level boost

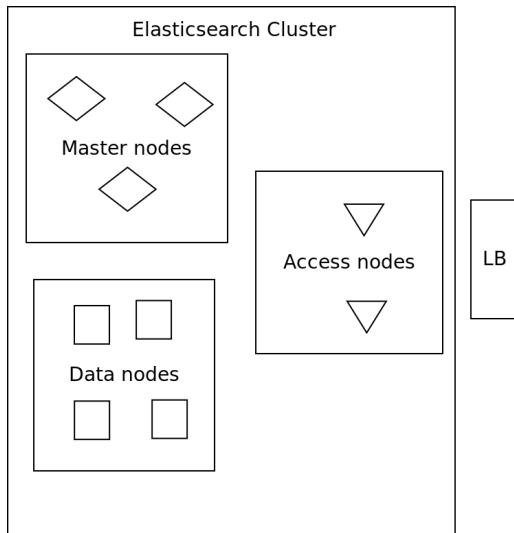
(<https://www.elastic.co/guide/en/elasticsearch/guide/current/practical-scoring-function.html>)

- 1 Master node - keeping, updating, broadcasting state of the cluster (list of nodes, mappings)
- 2 Data node - storing data and executing searches on shard
- 3 Access node - forwarding requests to data nodes and merging results
- 4 Ingest node - preprocess documents before indexing them

# Leader election

- 1 zen discovery
- 2 leader election
- 3 quorum

# Cluster



# Q&A



<https://github.com/jpodeszwik/elasticsearch-bottom-up>